

# Process and Thread

Process concept. Process management. PCB and context switch. Short-term vs. long-term CPU scheduling. Process creation and termination. Inter-process communication: shared memory and message passing. Thread vs. process. Kernel vs. user thread.

OS3: 1/2/2016

Textbook (SGG): Ch. 3.1-3.3,3.4.1,3.6.1,4.1-4.2,4.3.1,4.4,4.5.1-4.5.2



# Process Concept

---


- An operating system executes a variety of programs:
  - Batch system – jobs
  - Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably
- Process – a program in execution; process execution must progress in sequential fashion
  - Program is static (a file); process is dynamic
  - You can run the same program  $n$  times (e.g., edit  $n$  files): one program,  $n$  processes
- A process defines a line of *concurrency*, includes:
  - program counter
  - stack
  - data section





# Process in Memory

- Process also defines an address space
  - Address space is *private*
  - Not accessible (by default) from another process
- Hence, process couples **two** abstractions
  - Concurrency
  - Protection
- Can OS determine direction of stack growth?
- Advantage of stack and heap growing in opposite directions?





# Process State

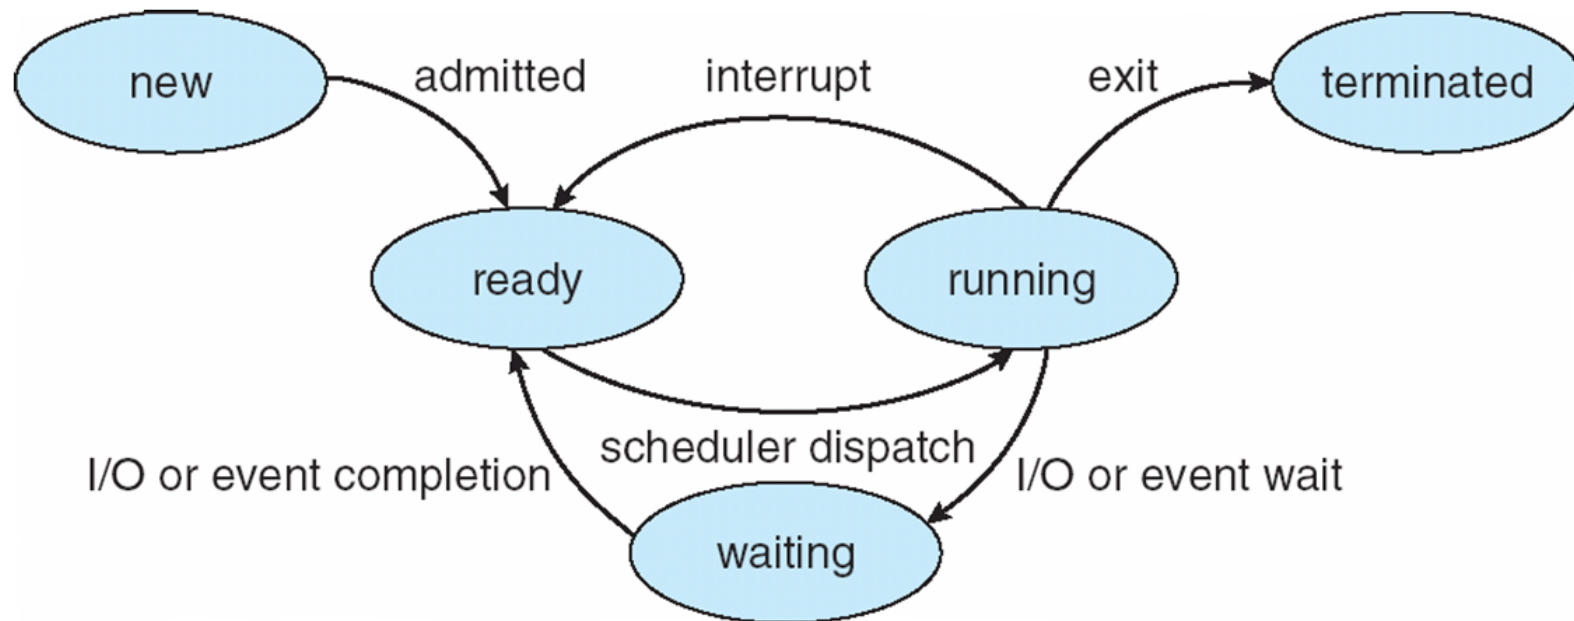
---

- As a process executes, it changes *state*
  - **new**: The process is being created
  - **running**: Instructions are being executed
  - **waiting** (or **blocked**): The process is waiting for some event to occur
  - **ready**: The process is waiting to be assigned to a processor (runnable)
    - What's needed to change it from runnable to running?
  - **terminated**: The process has finished execution





# Diagram of Process State





# Process Control Block (PCB)

---

To manage a process, OS keeps information about each process in a data structure called PCB

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information





# Process Control Block (PCB)

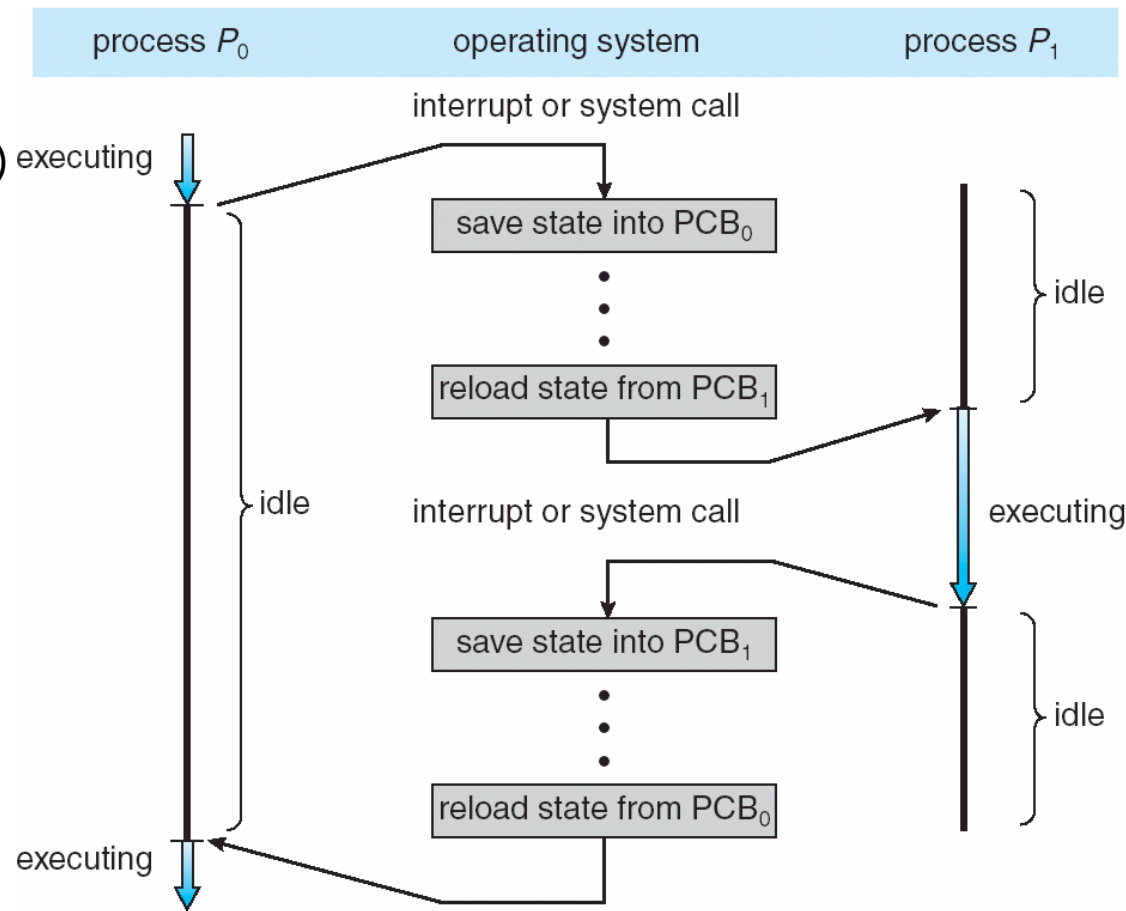
---





# CPU Switch From Process to Process

- Processes run independently (logically separate) and concurrently
- For a uniprocessor system, this is an illusion
  - As if each process has its own CPU
  - Can view it as *virtual CPU*







# Process Scheduling Queues

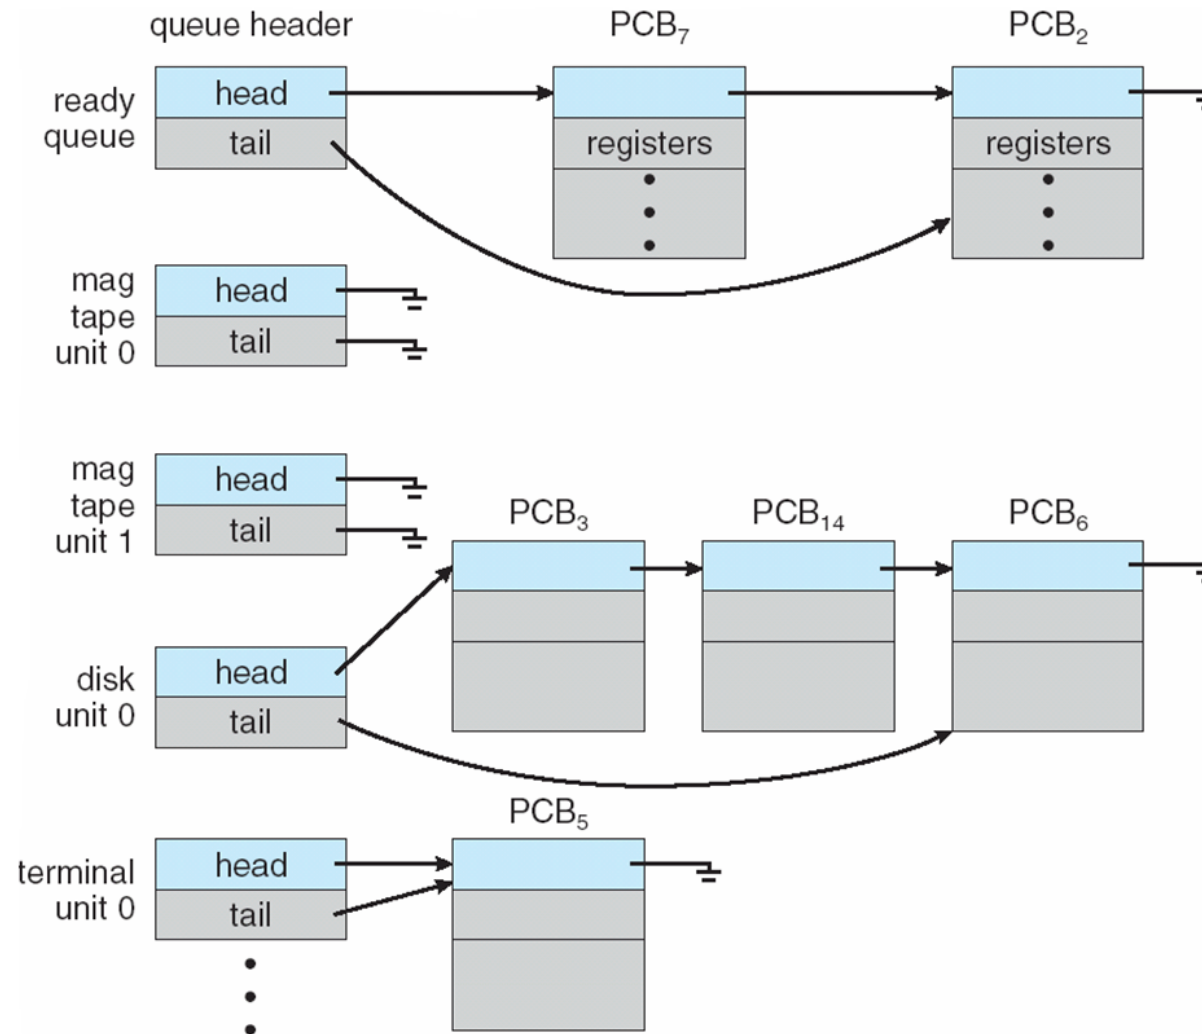
---

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Process migrates among the various queues during execution



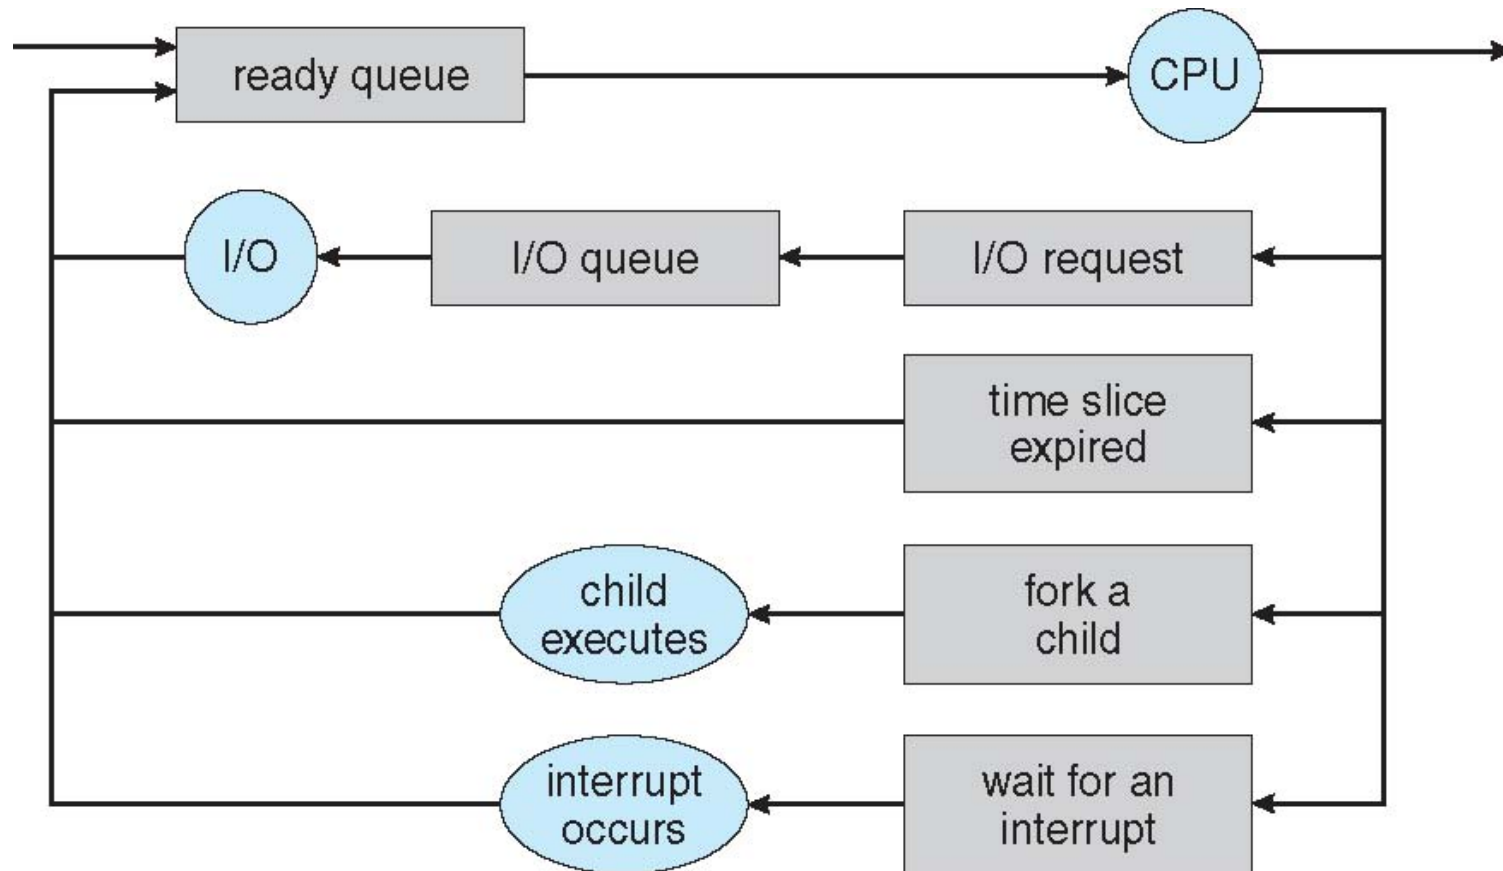


# Ready Queue And Various I/O Device Queues





# Representation of Process Scheduling





# Two Kinds of CPU Schedulers

---

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
  - i.e., loaded into main memory (swapped in)
  - If memory is limited, some processes can be *swapped out* (where are they then?)
- **Short-term scheduler** (or CPU scheduler) – selects which (in-memory) process should be executed next and allocates CPU to that process





## Schedulers (Cont.)

---

- Short-term scheduler is invoked very frequently (e.g., milliseconds) ⇒ must be fast
- Long-term scheduler is invoked very infrequently (seconds, minutes) ⇒ can be slow
- The long-term scheduler controls the *degree of multiprogramming*
  - i.e., How many processes are allowed to run at the same time
- Processes can behave quite differently, e.g.,
  - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - **CPU-bound process** – spends more time doing computations; few very long CPU bursts





# Context Switch

---

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**
- **Context** (where the process is in its execution) of a process stored in the PCB
- Context-switch time is *overhead*
  - System does no useful work while switching
  - So, don't want to do it too much, but what is *just enough*?
- Time dependent on hardware support





# Process Creation

---

- Processes form a family tree!
- **Parent** process create **child** processes, which in turn create other processes, forming a tree of processes
- Generally, process identified and managed via **process identifier (pid)**
- Parent/children can share resources (e.g., opened files) in different ways
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution
  - Parent and children execute concurrently
  - Parent can wait for children to terminate (**wait()** system call)





# Process Creation (Cont.)

---

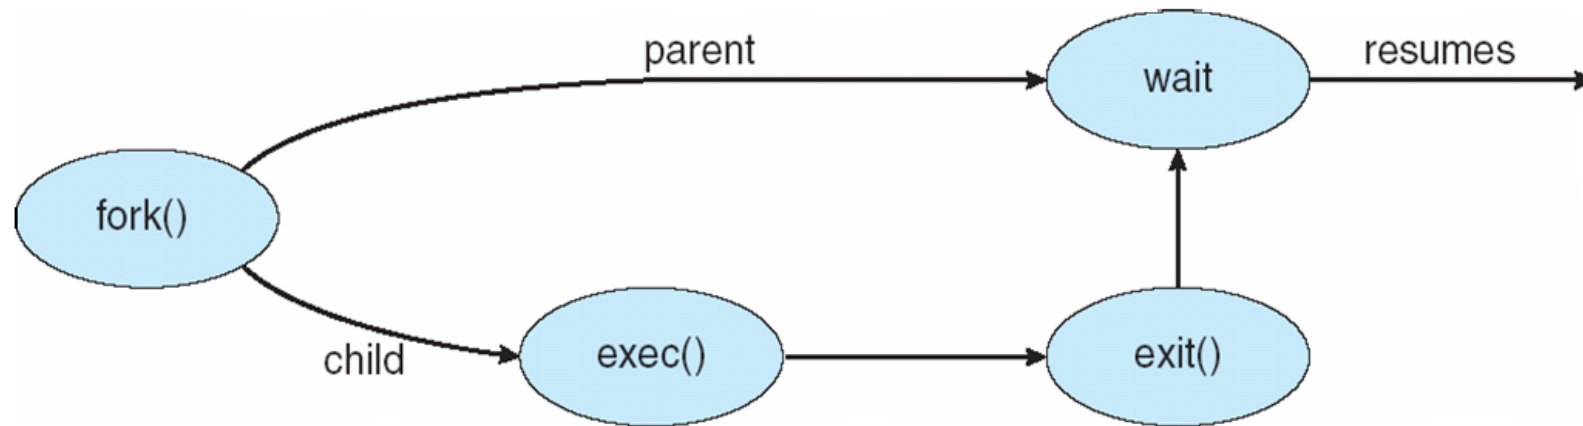
- Address space
  - Child gets its own address space, whose content is initially a duplicate of parent's
  - Child then usually loads a new program into its address space
- UNIX examples
  - **fork** system call creates new process
  - **exec** system call (used after a **fork**) loads new program into the process's memory space







# Process Creation



- After parent creates child, execution for both resumes as return from `fork()`
- How do you tell parent's return from child's return then?





# C Program Forking Separate Process

---

```
int main()
{
    pid_t  pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```





# Process Creation in Java

---

```
import java.io.*;

public class OSProcess
{
    public static void main(String[] args) throws IOException {
        if (args.length != 1) {
            System.err.println("Usage: java OSProcess <command>");
            System.exit(0);
        }

        // args[0] is the command
        ProcessBuilder pb = new ProcessBuilder(args[0]);
        Process proc = pb.start();

        // obtain the input stream
        InputStream is = proc.getInputStream();
        InputStreamReader isr = new InputStreamReader(is);
        BufferedReader br = new BufferedReader(isr);

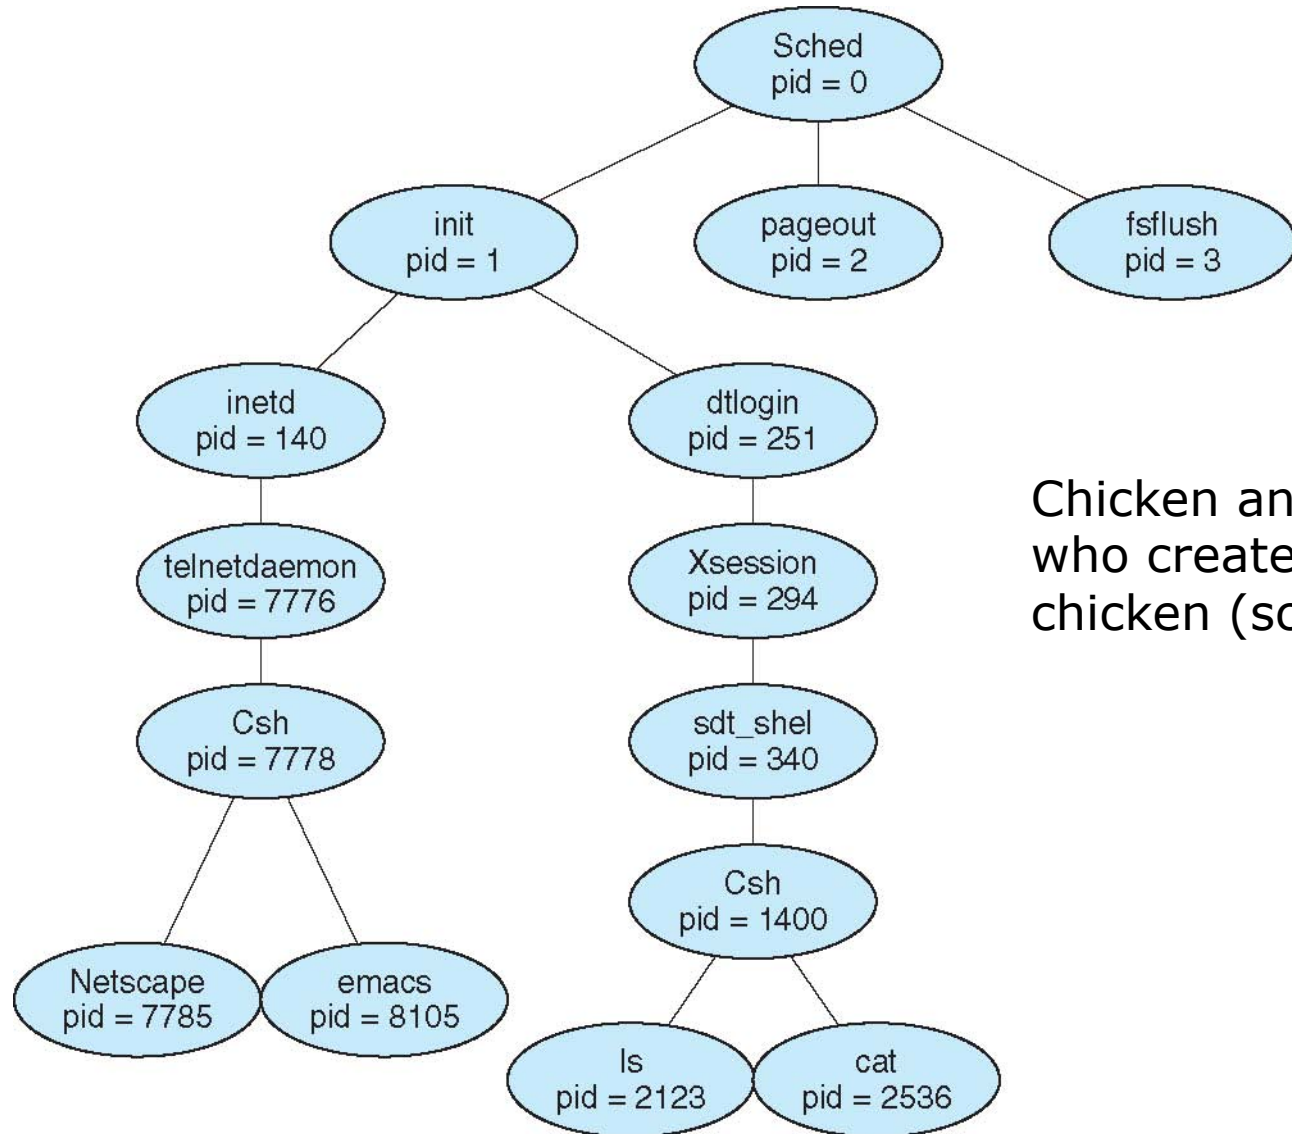
        // read what is returned by the command
        String line;
        while ( (line = br.readLine()) != null)
            System.out.println(line);

        br.close();
    }
}
```





# A tree of processes (Solaris OS)



Chicken and egg:  
who creates first  
chicken (sched)?





# Process Termination

---

- Process executes last statement and informs OS (via **exit()** system call)
  - Output data from child to parent (via **wait**)
  - Process's resources may be deallocated by OS
    - But sometimes, process can also still exist as “zombie”
- Parent may terminate execution of child processes (**abort**)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent exits
    - Some operating systems don't allow its children to continue
      - All children terminated - **cascading termination**
    - UNIX has *job control* feature that defines different behaviors
      - So a UNIX job means a related group of processes, not synonymous with process





# Interprocess Communication

---

- Processes are by default **independent**, but they can also agree to be **cooperating**
- Cooperating processes may affect each other, mainly through sharing data
- Reasons for cooperating processes:
  - Share information
  - Speed up computation
  - Achieve modularity (hence protection) in spite of cooperation
  - Convenience (e.g., “ls -l | wc -l” roughly counts how many files you have)
- Cooperating processes need **interprocess communication (IPC)**
- Two basic models of IPC
  - Shared memory
  - Message passing





## Activity 3.1: Multiprocess Program

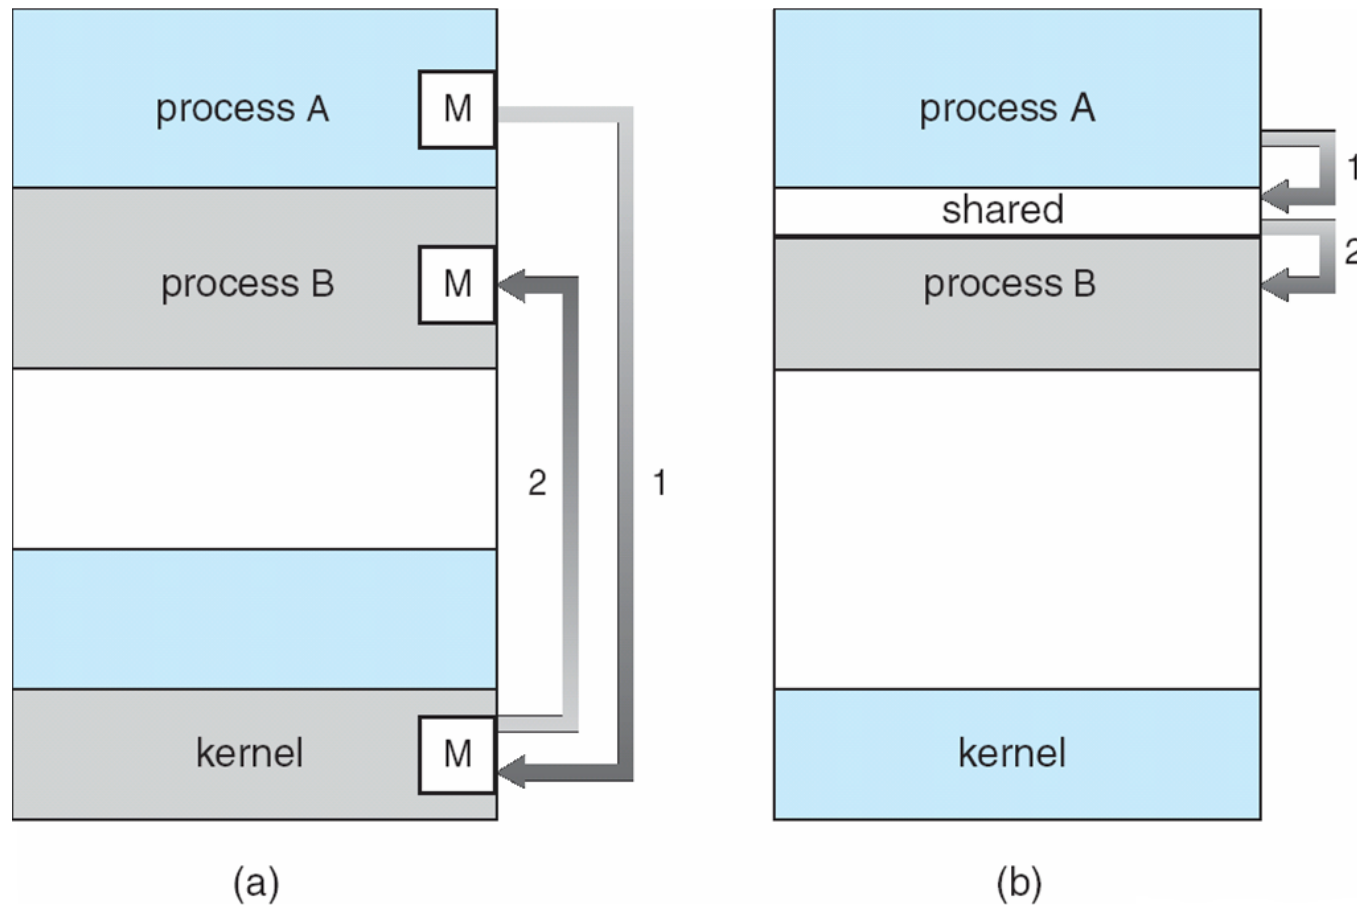
---

- If you run a task as  $N$  processes, what is the maximum speedup you can expect over a single-process implementation?
- If  $x\%$  of the task is sequential (i.e., can't be parallelized), what then is the maximum speedup?
- On a practical system, what other important factors will limit your actual speedup below the maximum possible?





# Communications Models







# Producer-Consumer Problem

---

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - Items produced and waiting to be consumed are put in *shared* buffer
  - *unbounded-buffer* places no practical limit on the size of the buffer
  - *bounded-buffer* assumes that there is a fixed buffer size





# POSIX Example of Shared Memory IPC

---

## ■ POSIX Shared Memory

- Process first creates shared memory segment (w/ read, write access for owner)

```
segment id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);
```

- Process wanting access to that shared memory must attach to it

```
shared_memory = (char *) shmat(id, NULL, 0);
```

- Now the process could write to the shared memory

```
sprintf(shared_memory, "Writing to shared memory");
```

- When done a process can detach the shared memory from its address space

```
shmdt(shared_memory);
```





# Message Passing Example: Sockets

---

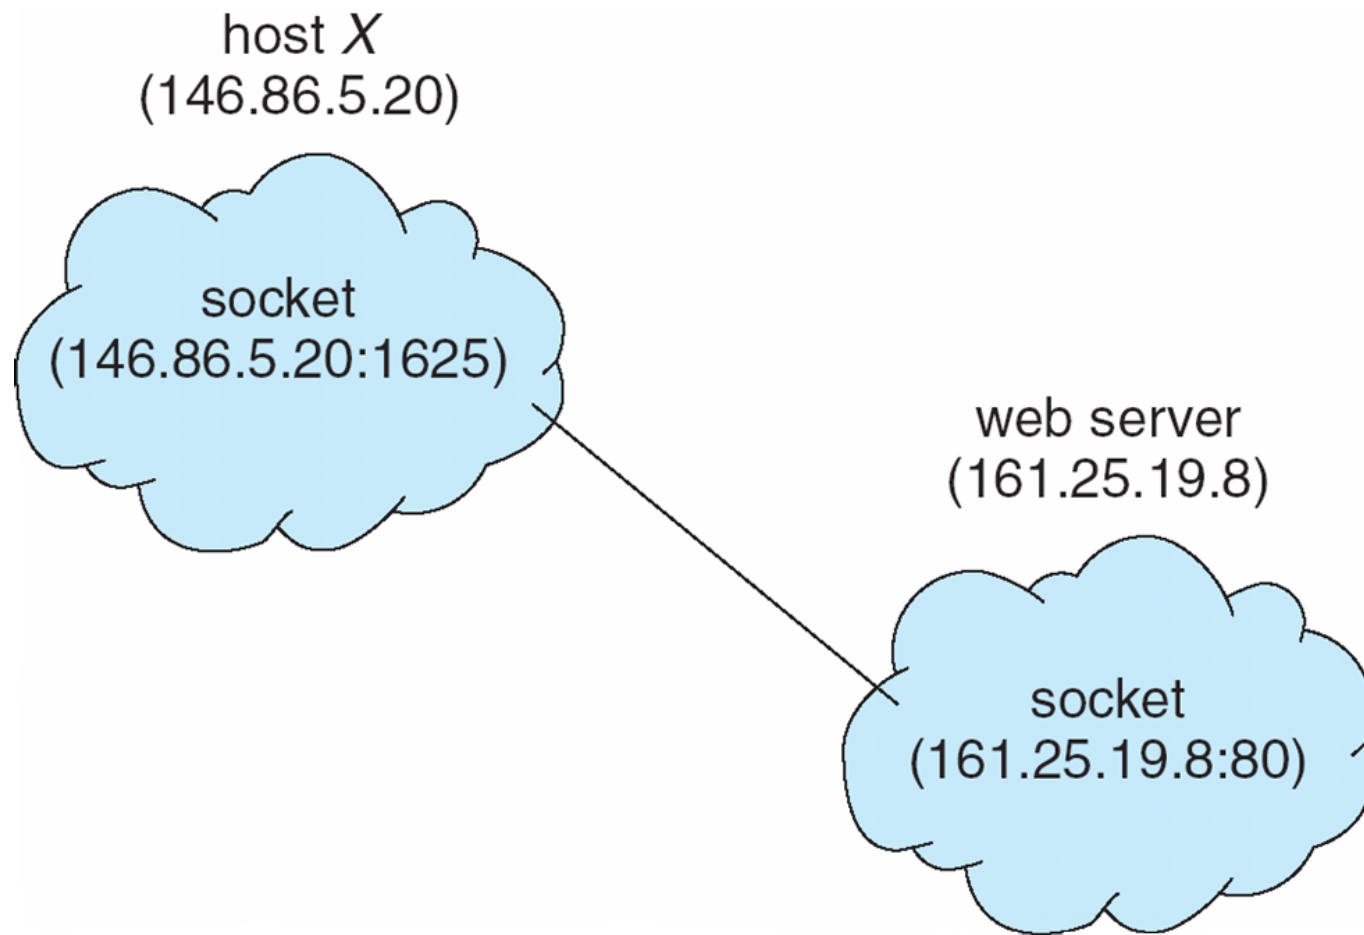
- A socket is defined as an *endpoint for communication*
  - Usually for network communication (e.g., TCP/IP), but also IPC within single machine
- Endpoint specified as concatenation of IP address and (TCP or UDP) port
  - E.g., **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication occurs between a pair of sockets – two flavors:
  - Connection oriented (e.g., TCP)
  - Connectionless (e.g., UDP)
- More details when we study networking





# Socket Communication

---





# Socket Communication in Java

---

```
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            // now listen for connections
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                // write the Date to the socket
                pout.println(new java.util.Date().toString());

                // close the socket and resume
                // listening for connections
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





# Socket Communication in Java

---

```
public class DateClient
{
    public static void main(String[] args) {
        try {
            //make connection to server socket
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            // read the date from the socket
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            // close the socket connection
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





# Thread vs. Process


---

- Recall: Process couples two abstractions: concurrency and protection
- Can I decouple the two, e.g., have concurrency *without* protection?
  - Yes, use *threads*
  - Many threads can run within a process, share the process's address space
    - ▶ No protection between them
    - ▶ But IPC is simple (e.g., no need for shmget, shmat, etc) and fast (much less to save/restore at context switch)

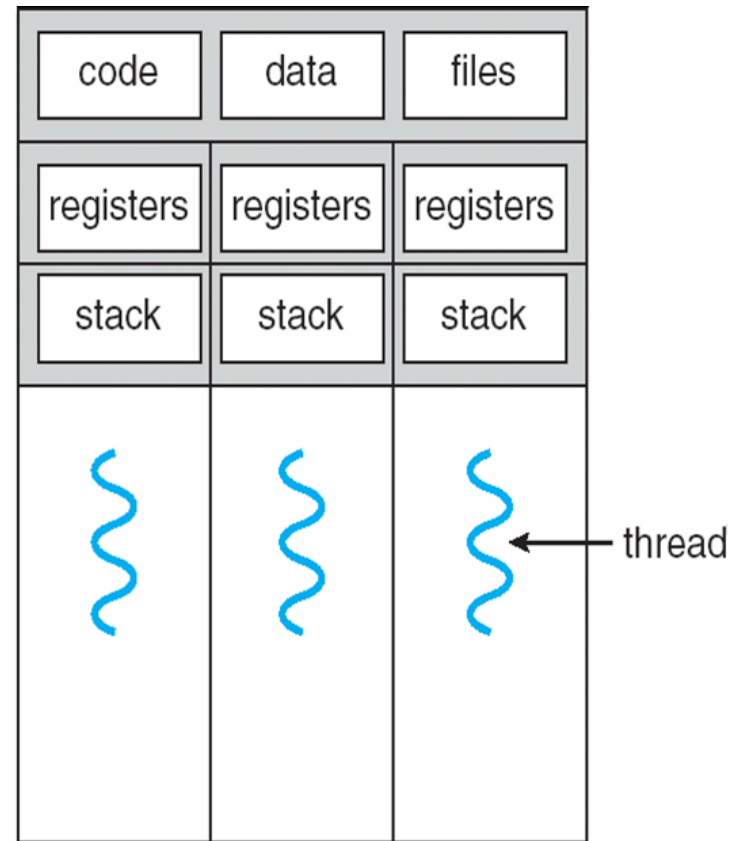




# Single and Multithreaded Processes



single-threaded process



multithreaded process







# Why (or why not) threads?

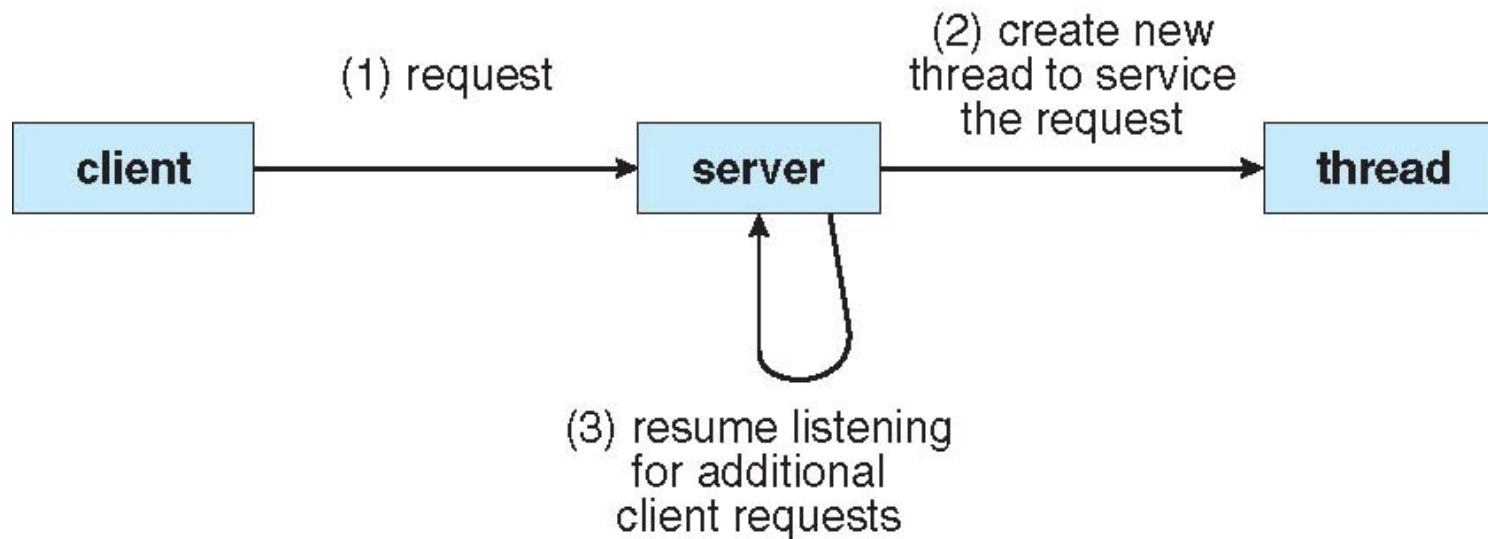
---

- Speedup by parallel execution (e.g., your Lab 2)
  - On multiprocessor or multicore systems
- Responsiveness
  - While one thread is blocked for IO, another thread can be executing and doing useful computation
- Logical modularity
  - Though without fault isolation
- Disadvantages
  - Context switch + synchronization overheads
  - Can be much harder to program and get right!





# Multithreaded Server Architecture





# Two types of threads

---

## ■ *Kernel* threads

- Known to OS kernel
- Scheduled by kernel CPU
- Take up kernel data structure (e.g., Thread Control Block like PCB)
- More expensive

## ■ *User* threads

- Not known to OS kernel
- Scheduled by thread scheduler (running in user mode) in thread library (e.g., POSIX pthread or Java threads) linked with process
- Less expensive





# Mapping from user to kernel threads

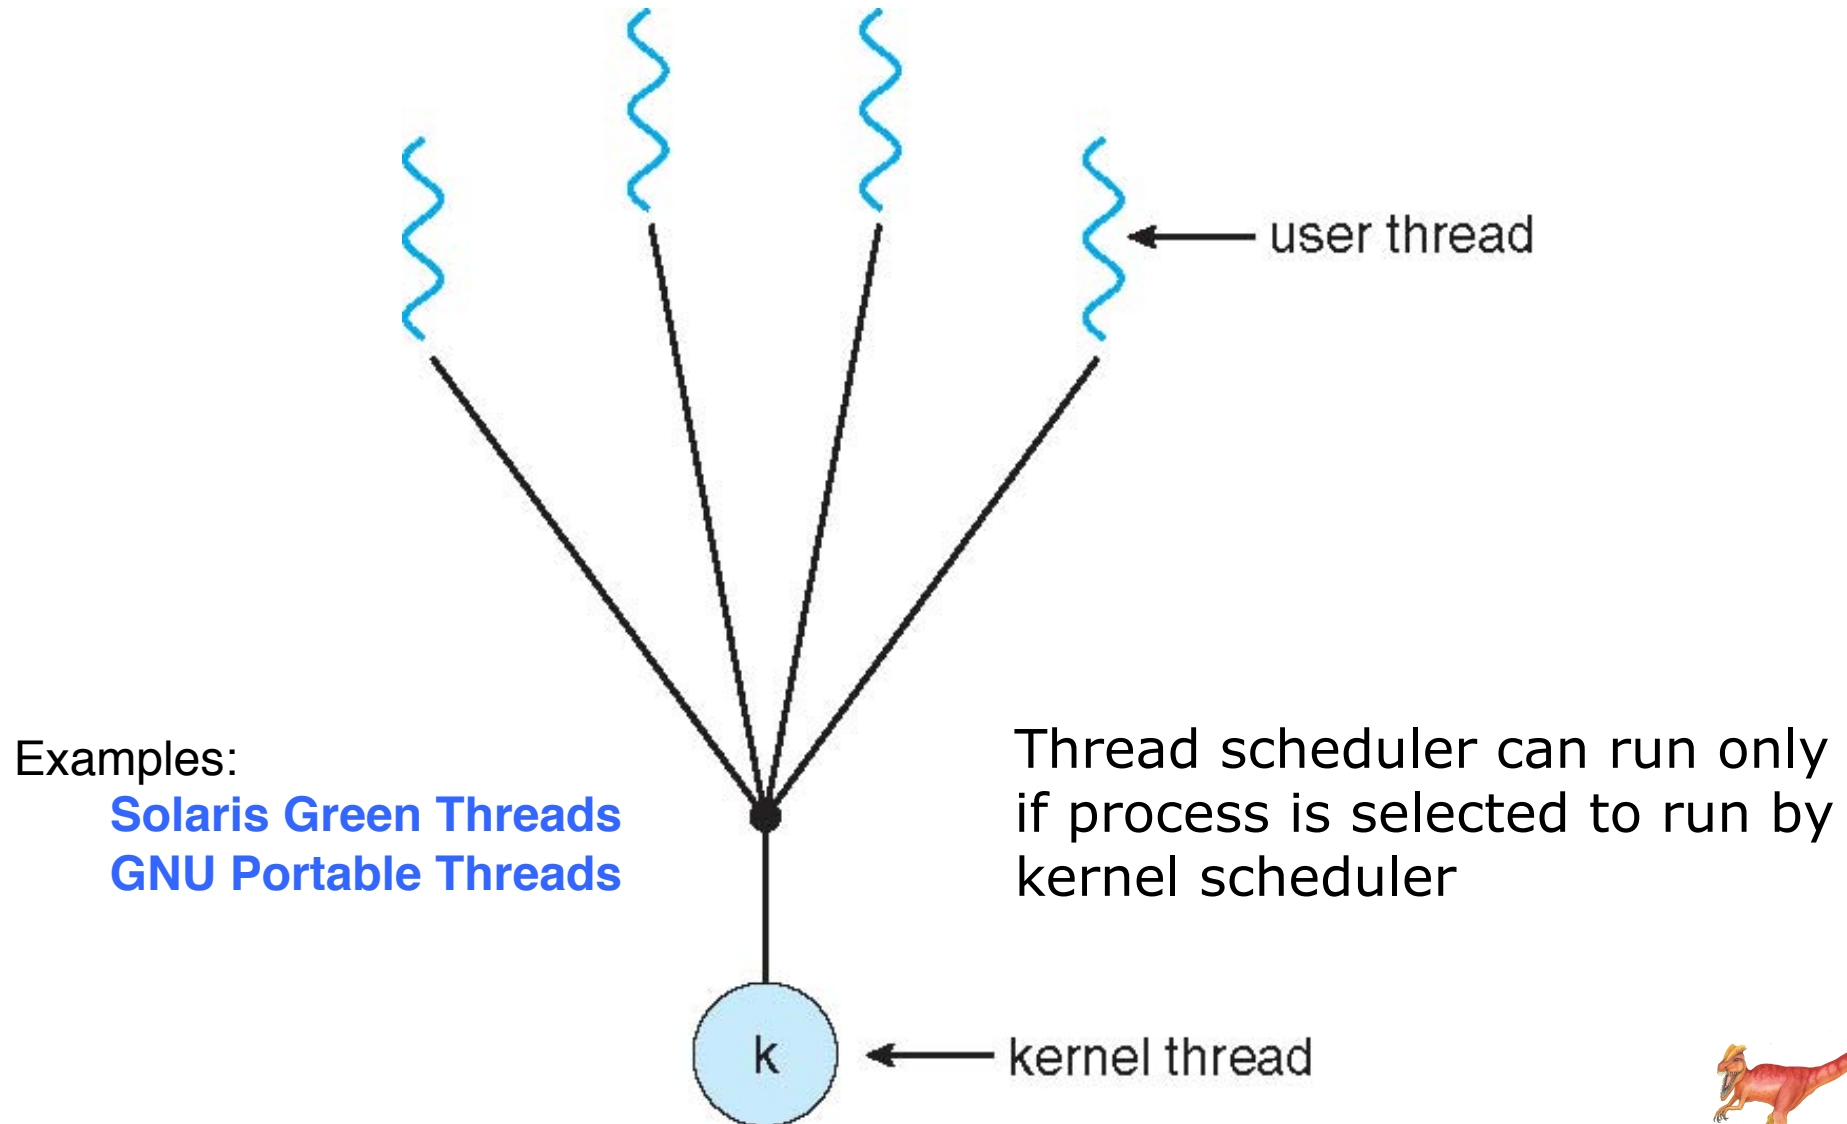
---

- Many-to-One
- One-to-One
- Many-to-Many
  - Two-level model






# Many-to-One Model





# One-to-one Model

---



## Examples

Windows NT/XP/2000

Linux

Solaris 9 and later






# Many-to-Many Model

## Examples

Solaris before v9 (Solaris LWP is user thread's door to kernel thread)


Windows NT/2000 w/ ThreadFiber package





# Two-level Model

Like many-to-many, except possible to bind user thread to kernel thread



Examples  
IRIX  
HP-UX  
Tru64 UNIX  
Solaris 8 and earlier







## Activity 3.2: Kernel vs. User Threads

---

- You have a multithreaded process in execution on a uniprocessor. One of the threads is in the middle of processing a previously accepted user command and it is runnable. Another thread, which is running, attempts to read a user command from the network and blocks in the kernel.
- What will happen next if these threads are kernel threads?
- What will happen next if these threads are user threads?
- Why is it generally faster to context switch user threads than kernel threads?





# Pthreads

---

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to developers of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS/X)





# Java Threads

---

- Java threads are managed by the JVM
- Java threads may be created by:
  - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```





# Java Threads - Example Program

---

```
class MutableInteger
{
    private int value;
    public int getValue() {
        return value;
    }
    public void setValue(int value) {
        this.value = value;
    }
}

class Summation implements Runnable
{
    private int upper;
    private MutableInteger sumValue;
    public Summation(int upper, MutableInteger sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }
    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setValue(sum);
    }
}
```





# Java Threads - Example Program

---

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                // create the object to be shared
                MutableInteger sum = new MutableInteger();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sum));
                thrd.start();
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sum.getValue());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>");
    }
}
```





# Java Thread States

